

JAVA INTERVIEW PREPARATION

CHAPTER 16

VIRTUAL THREADS AND THREAD-PER-REQUEST SCALABILITY

Senior / Lead Java Developer Textbook Edition

Full reading material - not summarized

Virtual Threads and Thread-per-Request Scalability: Cheap Concurrency with Real Resource Limits

Senior / Lead Java Developer Reading Chapter


Virtual threads became a permanent feature in Java 21. They are Java Thread instances scheduled primarily by the JVM rather than mapped one-to-one to long-lived operating-system threads. This makes it practical to create a new thread for each task and write straightforward blocking code while supporting far more concurrent waiting operations than a traditional platform-thread pool.

Virtual threads improve the cost model of threads; they do not make databases, CPUs, memory, sockets, or downstream services unlimited. Senior developers must separate concurrency from capacity, understand pinning and ThreadLocal implications, preserve cancellation, and measure the whole request path.

Platform Threads and Virtual Threads

A traditional platform thread is closely associated with an operating-system thread. OS threads have significant stack and scheduling cost, so applications commonly reuse a limited pool.

A virtual thread is mounted on a smaller set of carrier platform threads while running. When a supported blocking operation occurs, the JVM can unmount the virtual thread, freeing the carrier to run another virtual thread.



The virtual thread retains its logical call stack and resumes later. Application code can keep sequential control flow without occupying one OS thread for every wait.

Creating Virtual Threads

Start one directly:

```
Thread thread = Thread.startVirtualThread(() -> process(order));
thread.join();
```

Use a builder for naming and configuration:

```
Thread thread = Thread.ofVirtual()
    .name("order-", 0)
    .start(() -> process(order));
```

For task submission, Java 21 provides an executor that creates a new virtual thread per task:

```
try (ExecutorService executor =
    Executors.newVirtualThreadPerTaskExecutor()) {

    Future<Customer> customer = executor.submit(
        () -> customerClient.load(customerId)
    );
    Future<List<Order>> orders = executor.submit(
        () -> orderClient.loadFor(customerId)
    );

    return new CustomerPage(customer.get(), orders.get());
}
```

Closing the executor waits for submitted tasks to complete. In production, timeouts, cancellation, partial failure, and interruption still require an explicit policy.

Do Not Pool Virtual Threads

Platform threads are pooled because they are scarce. Virtual threads are intended to be created per task.

platform-thread model	virtual-thread model
-----	-----
limited reusable workers	new virtual thread per task
task waits in executor queue	task begins, then may park
pool size limits concurrency	explicit resource limit does

Putting virtual threads into a small fixed pool removes much of their benefit and confuses thread scarcity with resource protection.

If a database permits 50 concurrent queries, protect it with a connection pool or Semaphore rather than a pool of 50 virtual threads for every kind of work.

```
Semaphore databasePermits = new Semaphore(50);

Customer loadCustomer(CustomerId id) throws InterruptedException {
    databasePermits.acquire();
    try {
        return repository.load(id);
    } finally {
        databasePermits.release();
    }
}
```

The virtual-thread population can remain large while the constrained dependency receives controlled concurrency.

Throughput, Latency, and Little's Law

Virtual threads primarily improve throughput for workloads that spend substantial time waiting. They do not make one remote call faster.

concurrency approximately = throughput x average latency

If a service handles 2,000 requests per second and each request waits one second, it needs roughly 2,000 concurrent in-flight executions. Virtual threads make representing that waiting concurrency cheaper.

If the work is CPU-bound, the number of useful simultaneously running tasks remains near available cores. Ten thousand virtual threads performing compression do not create more CPU. They add scheduling and memory overhead.

The Thread-per-Request Model

The classic synchronous model becomes scalable again for many I/O-heavy services:

```
OrderView handle(OrderId id) {
    Order order = orderRepository.load(id);           // blocks
    Customer customer = customerClient.load(order.customerId());
    Price price = pricingClient.calculate(order);
    return new OrderView(order, customer, price);
}
```

Sequential code retains normal stack traces, exceptions, loops, and debugging. Virtual threads can reduce the need to express every wait through callbacks or reactive operators solely to avoid blocking platform threads.

Reactive systems still offer valuable streaming, backpressure, event processing, and ecosystem semantics. Virtual threads are not an instruction to rewrite every reactive system. Choose from workload, team expertise, integration libraries, and operational evidence.

Blocking Must Be Virtual-Thread Friendly

JDK blocking operations such as many socket and queue operations cooperate with virtual threads. Libraries that block through native code, long critical sections, or unsupported mechanisms may hold the carrier.

In Java 21, a virtual thread can be pinned to its carrier when it blocks while executing inside a `synchronized` method or block, or while executing native/foreign code. Pinning prevents that carrier from running another virtual thread during the wait.

```
synchronized void updateAndCallRemoteService() {
    updateState();
    remoteClient.call(); // blocking while monitor held: problematic in Java 21
}
```

The deeper design problem is holding a lock across remote I/O. Reduce the critical section:

```
void updateAndCallRemoteService() {
    Request request;
    synchronized (this) {
        request = updateAndCreateRequest();
    }
    remoteClient.call(request);
}
```

Later JDK releases improve some monitor-pinning behavior, but applications targeting Java 21 should understand its rules and should avoid blocking while holding locks regardless.

Pinning Is Not the Same as Blocking

Normal parking is expected and efficient. Pinning is important when it is frequent and long enough to exhaust carriers.

```
normal blocking
virtual thread parks -> carrier runs another task

pinned blocking
virtual thread waits + carrier also unavailable
```

Do not replace every synchronized block preemptively. Short, non-blocking critical sections are usually fine. Diagnose with Java Flight Recorder events, thread dumps, and the virtual-thread scheduler diagnostics

supported by the target JDK.

Fix long blocking operations under locks and obsolete library behavior rather than increasing carrier parallelism blindly.

Downstream Capacity and Backpressure

Cheap threads remove one natural bottleneck: the small executor queue. Without another control, the application may direct enormous concurrency at a database or remote service.

```
50,000 incoming tasks
    |
    v
50,000 virtual threads
    |
    v
database with 100 connections
    |
    v
long waits, timeouts, memory retention, retry amplification
```

Use bounded admission, connection pools, Semaphores, rate limits, timeouts, circuit breakers, and load shedding according to the dependency contract.

A virtual thread waiting for a connection is cheaper than a platform thread, but fifty thousand waiting request objects still retain memory and represent unacceptable latency.

ThreadLocal at Large Scale

Virtual threads support ThreadLocal. Existing request-context code can often continue working, but per-thread state now multiplies across potentially very many threads.

```
private static final ThreadLocal<byte[]> BUFFER =
    ThreadLocal.withInitial(() -> new byte[1024 * 1024]);
```

One megabyte per virtual thread is catastrophic at high concurrency. Review ThreadLocal values for size, cleanup, inheritance, and lifecycle. Avoid using ThreadLocal as an unbounded cache.

Scoped Values were preview features in the Java 21 era, not a permanent Java 21 feature. They are designed for bounded, immutable context sharing, but production use must follow the status and API of the actual target JDK.

Interruption and Cancellation

Virtual threads use the existing interruption model. Blocking methods may throw `InterruptedException`. Code must preserve cancellation:

```
try {
    return queue.take();
} catch (InterruptedException exception) {
    Thread.currentThread().interrupt();
    throw new TaskCancelledException(exception);
}
```

Sometimes the method can declare and propagate `InterruptedException` instead of translating it. Do not swallow interruption and continue lengthy work; shutdown, timeout, and parent cancellation may depend on it.

Cancelling a `Future` with interruption requests cancellation but cannot guarantee arbitrary code stops. Libraries, loops, and I/O operations must cooperate.

Timeouts Still Matter

Cheap waiting is still waiting. Every remote call, connection acquisition, and overall request needs a time budget.

```
HttpRequest request = HttpRequest.newBuilder(uri)
    .timeout(Duration.ofSeconds(2))
    .build();
```

Coordinate nested timeouts so children cannot outlive the parent request indefinitely. A database timeout, HTTP timeout, task timeout, and ingress deadline should form one intentional budget.

Retries require idempotency and backoff. Virtual threads can make it easy to create retry storms because each retry is cheap to represent.

Structured Concurrency Context

Structured concurrency treats related child tasks as one lexical unit: success, failure, cancellation, and joining are scoped together.

```
parent request
|
+-- load customer
+-- load orders
|
join results or cancel siblings on failure
```

Structured concurrency APIs were preview features in Java 21, while virtual threads were permanent. For Java 21 production code, preview use requires explicit compiler/runtime flags and accepts API evolution risk.

Even without preview APIs, adopt structured principles: track child Futures, define a deadline, cancel siblings when results are no longer needed, and do not leak background work beyond request ownership.

Framework Integration

A server or framework must deliberately execute requests on virtual threads. Merely running Java 21 does not automatically convert every platform-thread executor.

Check:

- Server and framework version support.
- Which request, scheduler, and client executors are virtual.
- JDBC and HTTP client blocking behavior.
- ThreadLocal security, tracing, and transaction context.
- Metrics that assume a small worker pool.
- Libraries using native calls or monitor-heavy synchronization.

Load test the complete stack. A virtual-thread capable servlet container does not make a limited JDBC pool or downstream API more scalable.

Observability and Thread Dumps

Traditional dashboards monitor a small thread pool's active count and queue depth. A thread-per-task design may have hundreds of thousands of virtual threads, so "thread count is high" is not itself an incident.

Observe:

- In-flight request count and age.
- Downstream permit and connection wait time.
- Timeout, cancellation, and rejection rates.
- Virtual thread creation and termination rates.
- Pinning events and carrier utilization.
- Allocation and retained request memory.
- CPU saturation and runnable task count.

Modern thread dumps and JFR can represent virtual-thread activity without relying on old assumptions. Name virtual threads or tasks meaningfully when operational tools can use the names, but avoid high-cardinality metric labels per thread.

CPU-Bound and Mixed Workloads

For CPU-bound algorithms, use bounded parallelism related to available processors. Virtual threads can still provide a clean task model, but an explicit CPU executor or parallel algorithm prevents excessive runnable competition.

Mixed services should isolate or limit expensive CPU stages:

```
Future<Report> report = cpuExecutor.submit(
    () -> renderLargeReport(data)
);
```

Do not assume the virtual-thread scheduler is an application-level resource partitioner. Define capacity around scarce work.

Migration Strategy

A safe adoption path is:

1. Identify I/O-heavy synchronous workloads constrained by platform-thread pools.
2. Upgrade libraries and framework versions.
3. Inventory ThreadLocal state and lock-held blocking.
4. Enable virtual threads for a bounded service path.
5. Add downstream Semaphores, timeouts, and admission controls.
6. Load test with production-like latency and failures.
7. Inspect pinning, memory, cancellation, and connection waits.
8. Expand gradually and compare operational behavior.

Do not begin by deleting every executor. Executors may encode isolation, priority, rate limits, scheduling, or CPU bounds beyond thread reuse.

Testing Virtual-Thread Code

Correctness tests should not depend on whether a task happens to run on a particular carrier. Test application contracts:

- Successful completion and failure propagation.
- Interrupt and cancellation behavior.
- Timeout budgets.
- Semaphore release in finally blocks.
- Lock boundaries and no blocking within critical sections.
- ThreadLocal cleanup and context propagation.
- Downstream concurrency limits.

- Shutdown with active tasks.

Load tests should inject realistic database latency, remote failures, connection-pool constraints, and memory pressure. A benchmark using only `Thread.sleep` proves thread scalability, not application scalability.

Production Perspective

Virtual-thread incidents come from removing a thread bottleneck without replacing its capacity control. A service opens too many database transactions. Connection-pool waits retain tens of thousands of requests. A one-megabyte `ThreadLocal` multiplies across virtual threads. Blocking under synchronized pins carriers on Java 21. Interruptions are swallowed and timed-out requests continue writing. Dashboards alert on high thread count but miss downstream saturation.

When diagnosing:

- Is the workload waiting or CPU-bound?
- How many tasks are in flight, runnable, parked, and pinned?
- Which resource actually limits throughput?
- Are concurrency permits and queues bounded?
- What `ThreadLocal` state is retained per task?
- Are timeouts and cancellation propagated?
- Are locks held across blocking operations?
- Which framework executors really use virtual threads?
- What memory remains reachable from waiting tasks?

Interview-Ready Summary

Virtual threads are Java `Thread` instances scheduled by the JVM and became permanent in Java 21. They make thread-per-task and thread-per-request code practical for high-concurrency workloads that spend substantial time waiting.

Virtual threads improve throughput potential, not individual operation latency, CPU capacity, or downstream limits. Do not pool them to control concurrency; create one per task and use Semaphores, connection pools, rate limits, admission control, and timeouts around scarce resources.

Supported blocking usually parks and unmounts a virtual thread. In Java 21, blocking while holding a monitor or in native/foreign code can pin the carrier. Diagnose meaningful long pinning rather than replacing every synchronized block.

`ThreadLocal` works but can multiply memory across enormous thread counts. Interruption, cancellation, deadlines, retries, idempotency, transactions, and context propagation remain application responsibilities. Structured concurrency was preview in Java 21 and should not be confused with permanent virtual threads.

Revision Notes

- Virtual threads became permanent in Java 21.
- They are Java `Thread` instances scheduled primarily by the JVM.
- Virtual threads mount on carrier platform threads while running.
- Supported blocking can unmount a virtual thread and free its carrier.
- Create a virtual thread per task; do not use a small pool of virtual threads.
- `newVirtualThreadPerTaskExecutor` creates one virtual thread for each task.
- Virtual threads improve throughput for waiting workloads, not single-call latency.

- CPU-bound throughput remains limited by available cores.
- Little's Law relates concurrency, throughput, and latency.
- Use explicit limits around databases, remote APIs, and other scarce resources.
- A Semaphore or connection pool protects capacity better than pooling virtual threads.
- Cheap waiting can still retain unacceptable memory and latency.
- In Java 21, monitor-held blocking and native/foreign execution can pin carriers.
- Normal parking is not pinning.
- Avoid blocking I/O while holding locks.
- Diagnose pinning with JFR and target-JDK diagnostics.
- ThreadLocal state multiplies across potentially huge thread populations.
- Never use large per-thread buffers as an unbounded ThreadLocal cache.
- Preserve interruption and cancellation semantics.
- Cancelling a Future requires task and library cooperation.
- Every blocking dependency still needs a timeout.
- Coordinate nested operations within an overall request deadline.
- Retries still require idempotency, limits, and backoff.
- Structured concurrency was preview, not permanent, in Java 21.
- Frameworks must explicitly execute work on virtual threads.
- Verify JDBC, HTTP, ORM, tracing, security, and transaction behavior.
- High virtual-thread count alone is not an incident.
- Observe in-flight age, resource waits, pinning, CPU, memory, and cancellation.
- Isolate or bound CPU-heavy stages separately.
- Migrate incrementally and load test realistic latency and failures.
- Virtual threads remove thread scarcity, not system capacity limits.